

Logger for Cloud connector Go

Hello Everyone! In this page we would try to understand zap (Blazing fast, structured, leveled logging package for Go, see how it compares with other famous packages [here](#)) and how it is implemented at Avi.

As a prerequisites please make sure you have the following packages:-

- zap `go get -u go.uber.org/zap`
- lumberjack `go get gopkg.in/natefinch/lumberjack.v2`

We'll learn it all with an example!

Example

```
package main

import (
    "log"

    "go.uber.org/zap"
    "go.uber.org/zap/zapcore"
    "gopkg.in/natefinch/lumberjack.v2"
)

func logger1() {
    // In contexts where performance is nice, but not critical, use
    // the SugaredLogger. It's 4-10x faster than other structured logging
    // packages and supports both structured and printf-style logging.
    // SugaredLogger's structured logging APIs are loosely typed and accept a
    // variadic number of key-value pairs.
    logger, err := zap.NewDevelopment()
    if err != nil {
        log.Fatalf("can't initialize zap logger: %v", err) //
        // Replace this line with something as we don't want our program to stop
    }
    sugar := logger.Sugar() // Thus we see that it is easy to convert
    // to a SugaredLogger, do sugar.Desugar() to get back normal logger
    defer sugar.Sync()      // By default, loggers are unbuffered.
    // However, since zap's low-level APIs allow buffering, calling Sync before
    // letting your process exit is a good habit.

    // Info uses fmt.Sprint to construct and log a message
    // Infof uses fmt.Sprintf to log a templated message.
    // Infow logs a message with some additional context. The variadic
    // key-value pairs are treated as they are in With.
    // Read more about all this here:
    // https://godoc.org/go.uber.org/zap#SugaredLogger.Infow
    sugar.Infow("This is an INFO message with fields",
        "region", "us-west", // key-value pair 1
        "id", 2, // key-value pair 2
    )
}
```

```

    )
    // In the rare contexts where every microsecond and every
    allocation matter, use the Logger. It's even faster than the SugaredLogger
    and allocates far less, but it only supports strongly-typed, structured
    logging.
    // As mentioned simply before, logger doesn't have Infof nor Infow
    logger.Info("This is an INFO message with fields",
zap.String("region", "us-west"), zap.Int("id", 2))
}

// expected output:
// 2019-05-13T16:55:05.141+0530      INFO      tutorial/prodex1.go:22  This is
an INFO message with fields      {"region": "us-west", "id": 2}
// 2019-05-13T16:55:05.141+0530      INFO      tutorial/prodex1.go:22  This is
an INFO message with fields      {"region": "us-west", "id": 2}
// Also I wrote NewDevelopment (as this is the style we will mostly stick
with) but there are as well NewProduction and NewExample

func logger2() { // Custom Logger (This way I understand to be the best
among others)
    atom := zap.NewAtomicLevel()
    cfg := zap.Config{
        Encoding: "console", // Other setting is json
        Level:     atom,      // very important concept!
        // We have four levels, viz. Debug, Info, Warn, Error,
        Fatal. Putting DebugLevel here will imply that we can log all of these 5
        levels, if I were to put ErrorLevel, then I would only be able to log
        Error, Fatal level. This setting can be dynamically changed using our atom
        as shown below!

        // Note: Logging at Fatal will cause the program to
        terminate.

        OutputPaths:      []string{"stdout",
"/Users/sourabhaggarwal/Desktop/some.log"}, // can put multiple paths!
        ErrorOutputPaths: []string{"stderr"},
        InitialFields: map[string]interface{}{ // Initial fields
            "some": "bar"},
        EncoderConfig: zapcore.EncoderConfig{
            MessageKey: "message",

            LevelKey:     "level",
            EncodeLevel: zapcore.CapitalLevelEncoder,

            TimeKey:      "time",
            EncodeTime: zapcore.ISO8601TimeEncoder,

            CallerKey:     "caller",
            EncodeCaller: zapcore.ShortCallerEncoder,

            LineEnding: "\n-----*-----\n",
            // Although I didn't enable but we can as well
            have stack trace!
        },
    }
    atom.SetLevel(zap.DebugLevel)
}

```

```

    logger, _ := cfg.Build()
    defer logger.Sync()
    logger.Debug("Debug!")
    logger.Warn("Info will get ignored")
    atom.SetLevel(zap.WarnLevel) // This will make sure that info and
below doesn't get logged.
    logger.Info("This is an INFO message with fields",
zap.String("region", "us-west"), zap.Int("id", 2))
}

// Output:
// 2019-05-13T17:19:41.839+0530    DEBUG    tutorial/prodex1.go:65  Debug!
{"some": "bar"}
// -----
// 2019-05-13T17:19:41.839+0530    WARN     tutorial/prodex1.go:66  Info
will get ignored  {"some": "bar"}
// -----

func logger3() { // Now lets add lumberjack (helps in log rotation)
    atom := zap.NewAtomicLevel()
    atom.SetLevel(zap.DebugLevel)
    w := zapcore.AddSync(&lumberjack.Logger{
        // put you desired file path here!
        Filename: "/Users/sourabhaggarwal/Desktop/logName",
        // MaxAge is the maximum number of days to retain old log
files based on the
        // timestamp encoded in their filename. Note that a day
is defined as 24
        // hours and may not exactly correspond to calendar days
due to daylight
        // savings, leap seconds, etc. The default is not to
remove old log files
        // based on age.
        MaxAge: 28, // days
        MaxSize: 20, // megabytes
        // MaxBackups is the maximum number of old log files to
retain. The default
        // is to retain all old log files (though MaxAge may still
cause them to get
        // deleted.)
        MaxBackups: 3,
        Compress: true, // Yes we want to compress the files
    })
    core := zapcore.NewCore(
        // could have used here zap.NewDevelopmentEncoderConfig()
but still not getting line numbers...
        zapcore.NewConsoleEncoder(zapcore.EncoderConfig{ // same
old settings available
            MessageKey: "message",

            LevelKey: "level",
            EncodeLevel: zapcore.CapitalLevelEncoder,

            TimeKey: "time",

```

```

        EncodeTime: zapcore.ISO8601TimeEncoder,

        CallerKey:    "caller",
        EncodeCaller: zapcore.ShortCallerEncoder,

        LineEnding: "\n-----*-----\n",
    }},
    w,
    atom,
)
logger := zap.New(core, zap.AddCaller())
defer logger.Sync()
sugar := logger.Sugar()
sugar.Infow("This is an INFO message with fields", "region", "us-
west", "id", 2)
// logger.Info("This is an INFO message with fields",
zap.String("region", "us-west"), zap.Int("id", 2))
}

// Output is in your mentioned file!

func main() {
    logger1()
    logger2()
    logger3()
}

```

Implementation at Avi

- It's almost the same as *logger3* just that as we have wrapped everything inside a package, thus, it is therefore required to modify

```
logger := zap.New(core, zap.AddCaller())
```

to

```
logger := zap.New(core, zap.AddCaller(), zap.AddCallerSkip(1))
```

So as to show *that* caller which is invoking the suitable method of our package.

- Due to simplicity of *SugaredLogger* (as evident in *logger1* where its details are as well explained), we have decided to stick with it.

To see it head over to [avi-dev/go/src/avi/infra/avilog.go](https://github.com/avi-dev/go/src/avi/infra/avilog.go)